

# Improving Reliability Through Analyzing and Debugging Floating-Point Software

**Ignacio Laguna**

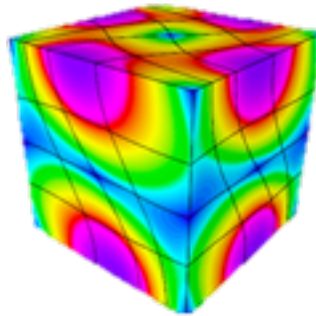
Computer Scientist  
Center for Applied Scientific Computing

2020 ECP Annual Meeting, Feb 4, 2020



# A Hard-To-Debug Case

Hydrodynamics mini application



Early development and porting to new system (IBM Power8, NVIDIA GPUs)

clang -O1:  $|e| = 129941.1064990107$   
clang -O2:  $|e| = 129941.1064990107$   
clang -O3:  $|e| = 129941.1064990107$

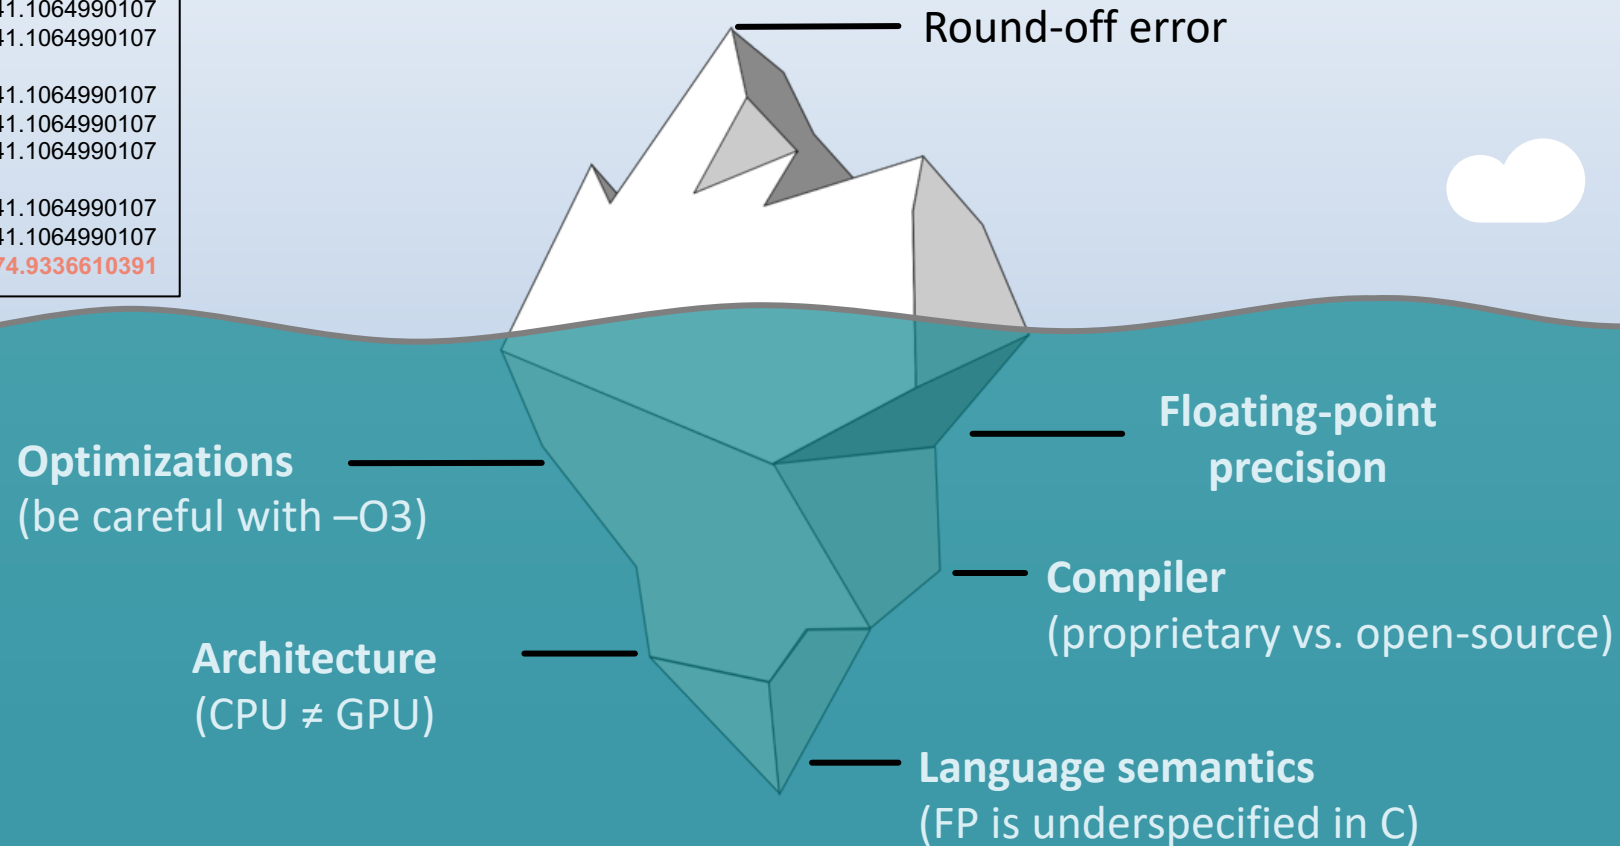
gcc -O1:  $|e| = 129941.1064990107$   
gcc -O2:  $|e| = 129941.1064990107$   
gcc -O3:  $|e| = 129941.1064990107$

xlc -O1:  $|e| = 129941.1064990107$   
xlc -O2:  $|e| = 129941.1064990107$   
xlc -O3:  $|e| = 144174.9336610391$

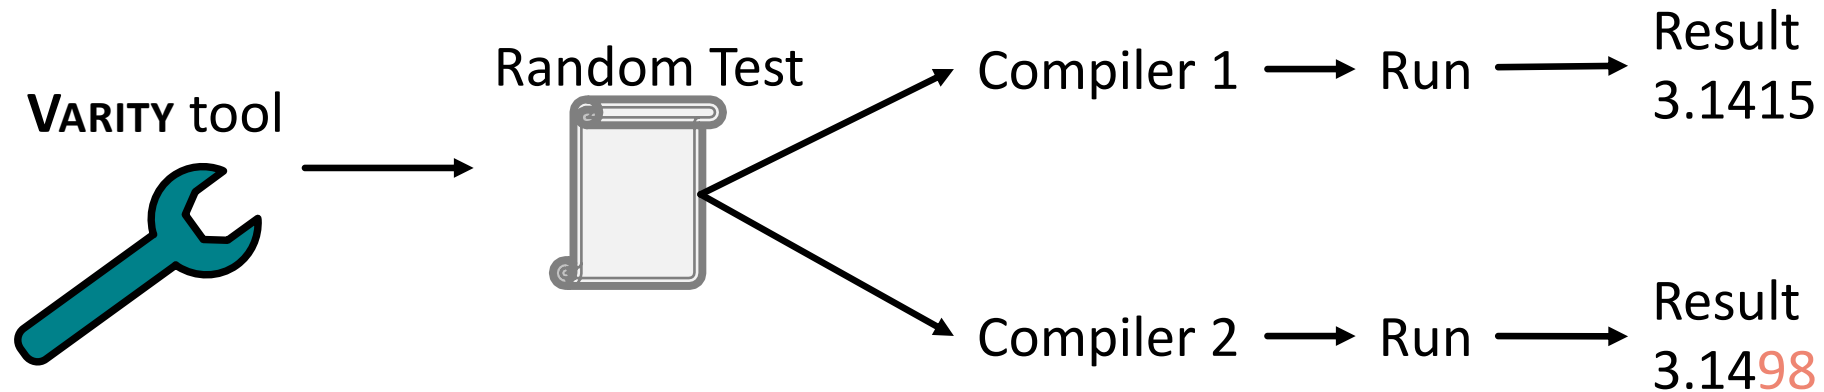
It took several weeks of effort to debug it

# Many Factors are Involved in Unexpected Numerical Results

clang	-O1:	e  = 129941.1064990107
clang	-O2:	e  = 129941.1064990107
clang	-O3:	e  = 129941.1064990107
gcc	-O1:	e  = 129941.1064990107
gcc	-O2:	e  = 129941.1064990107
gcc	-O3:	e  = 129941.1064990107
xlc	-O1:	e  = 129941.1064990107
xlc	-O2:	e  = 129941.1064990107
xlc	-O3:	e  = 144174.9336610391



# What Floating-Point Code Can Produce Variability?



<https://github.com/LLNL/Varity>

# Variability Examples Found by Varsity

## Example 1: variability between host and device

```
void compute(double comp,int var_1,double var_2,
double var_3,double var_4,double var_5,double var_6,
double var_7,double var_8,double var_9,double var_10,
double var_11,double var_12,double var_13,
double var_14) {
    double tmp_1 = +1.7948E-306;
    comp = tmp_1 + +1.2280E305 - var_2 +
        ceil((+1.0525E-307 - var_3 / var_4 / var_5));
    for (int i=0; i < var_1; ++i) {
        comp += (var_6 * (var_7 - var_8 - var_9));
    }
    if (comp > var_10 * var_11) {
        comp = (-1.7924E-320 - (+0.0 / (var_12/var_13)));
        comp += (var_14 * (+0.0 - -1.4541E-306));
    }
    printf("%.17g\n", comp);
}
```

### Input

```
0.0 5 -0.0 -1.3121E-306 +1.9332E-313 +1.0351E-306
+1.1275E172 -1.7335E113 +1.2916E306 +1.9142E-319
+1.1877E-306 +1.2973E-101 +1.0607E-181 -1.9621E-306
-1.5913E118-03
```

### clang -O3

```
$ ./test-clang
NaN
```

### nvcc -O3 (V100 GPU)

```
$ ./test-nvcc
-2.3139093300000002e-188
```

## Example 2: variability even with -O0

```
void compute(double tmp_1, double tmp_2, double tmp_3,
double tmp_4, double tmp_5, double tmp_6) {
    if (tmp_1 > (-1.9275E54 * tmp_2 + (tmp_3 - tmp_4 * tmp_5)))
    {
        tmp_1 = (0 * tmp_6);
    }
    printf("%.17g\n", tmp_1);
    return 0;
}
```

### Input

```
+1.3438E306 -1.8226E305 +1.4310E306 -1.8556E305 -
1.2631E305 -1.0353E3
```

### clang -O0

```
$ ./test-clang
1.3437999999999999e+306
```

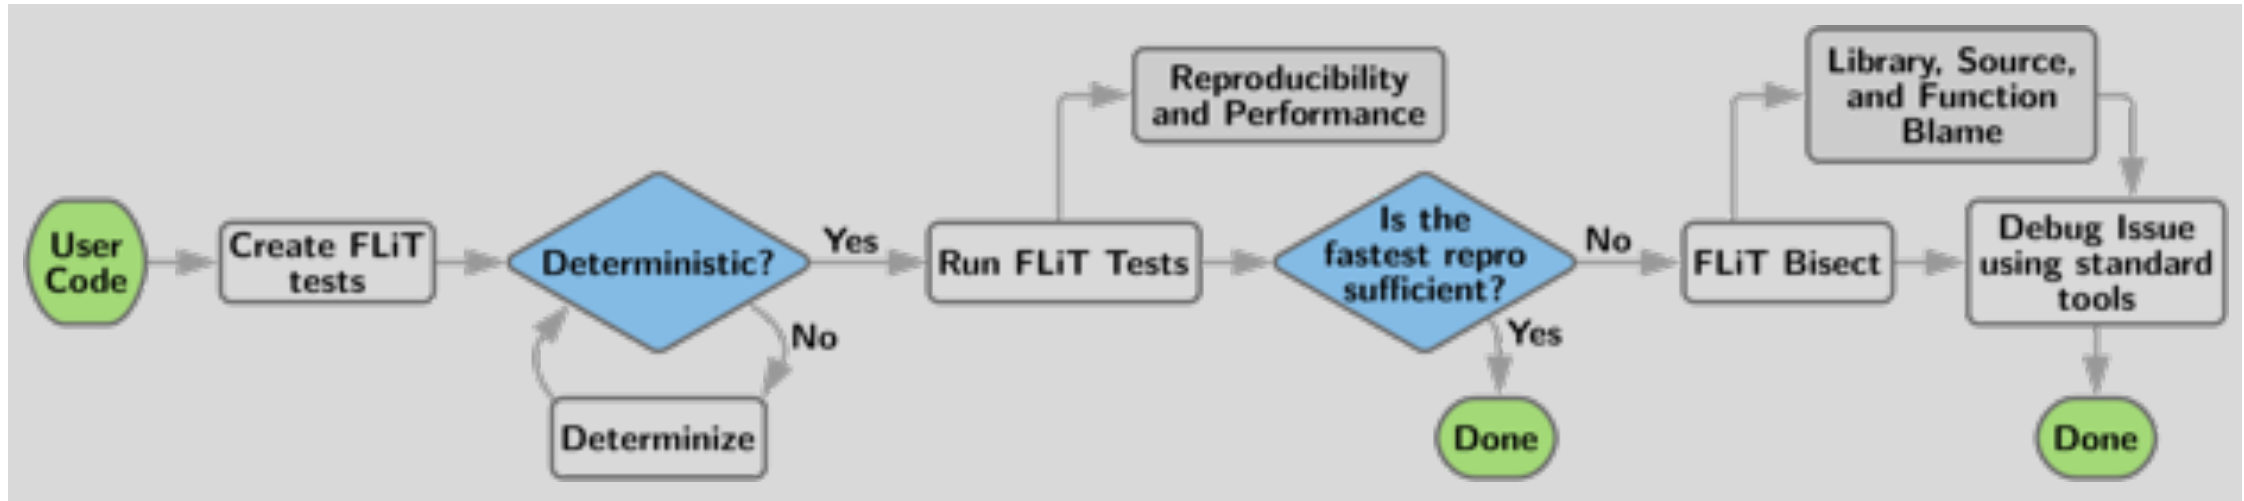
### gcc -O0

```
$ ./test-gcc
1.3437999999999999e+306
```

### xlc -O0

```
$ ./test-xlc
-0
```

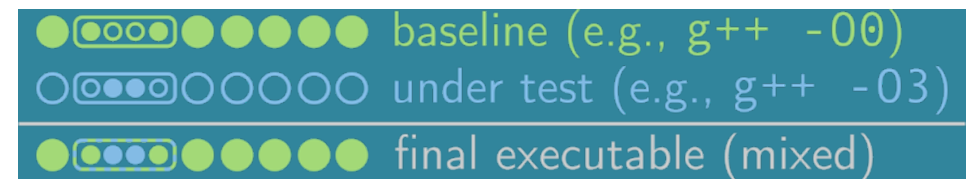
# FLiT: Floating-point Litmus Tester



## Multiple Levels:

- Determine variability-inducing compilations
- Analyze the tradeoff of reproducibility and performance
- Locate variability by identifying files and functions causing variability

## Bisection Method



Michael Bentley, Ian Briggs, Ganesh Gopalakrishnan, Dong H. Ahn, Ignacio Laguna, Gregory L. Lee, and Holger E. Jones. **Multi-Level Analysis of Compiler-Induced Variability and Performance Tradeoffs**. In Proceedings of the 28th International Symposium on High-Performance Parallel and Distributed Computing (HPDC '19).

# Detecting the Result of Exceptions in a CUDA Program

- Place **printf** statements in the code (as many as possible)

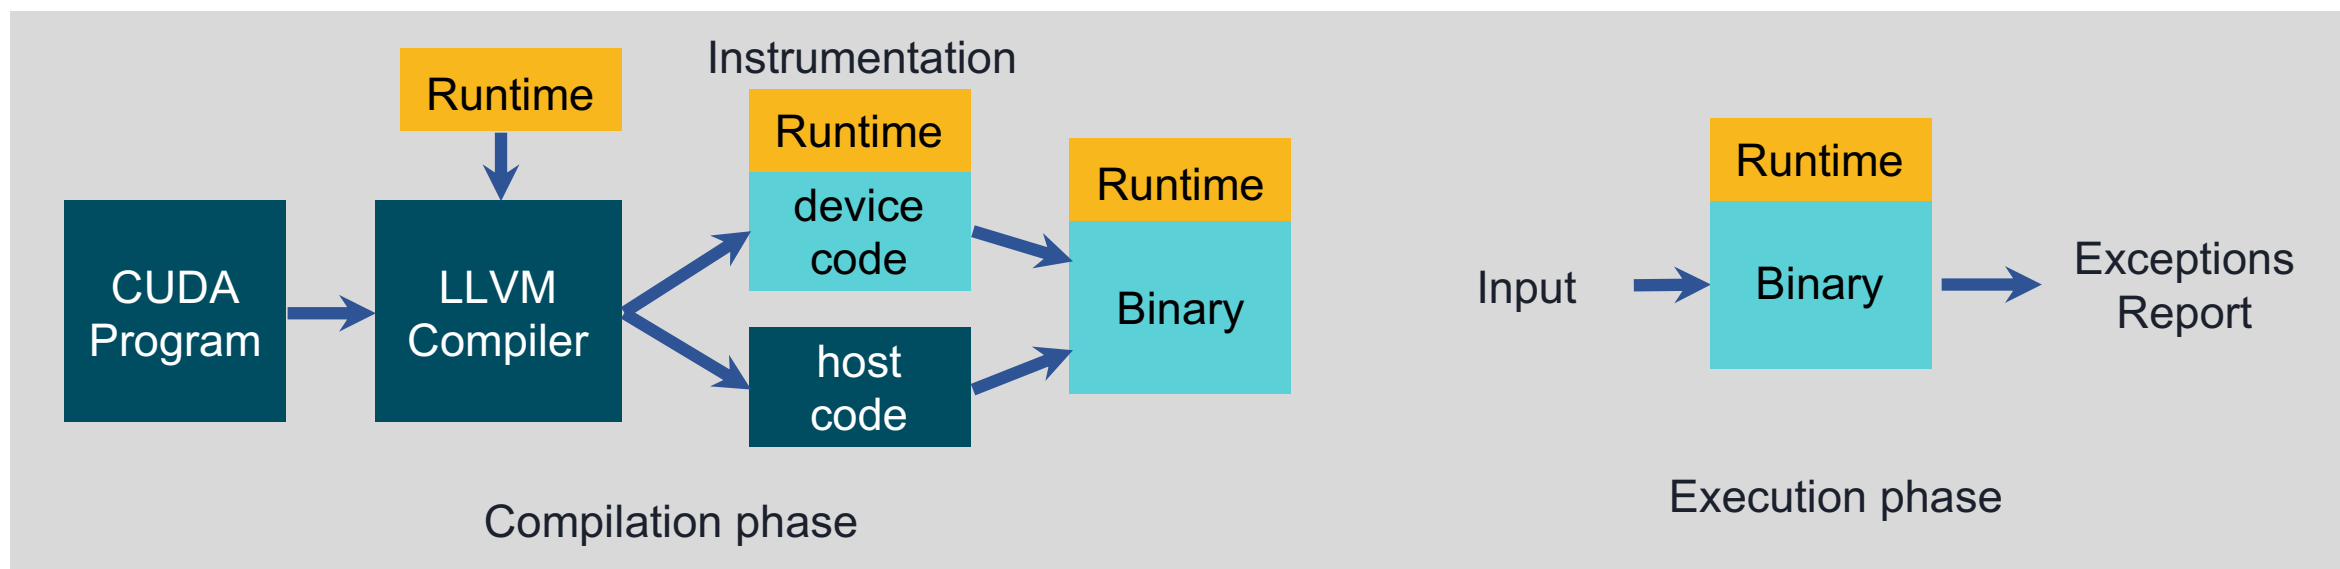
```
double x = 0;  
x = x/x;  
printf("res = %e\n", x);
```

- Programing checks are available in CUDA:

```
__device__ int isnan ( float  a );  
__device__ int isnan ( double a );
```

These solutions are not ideal; they require significant programming effort

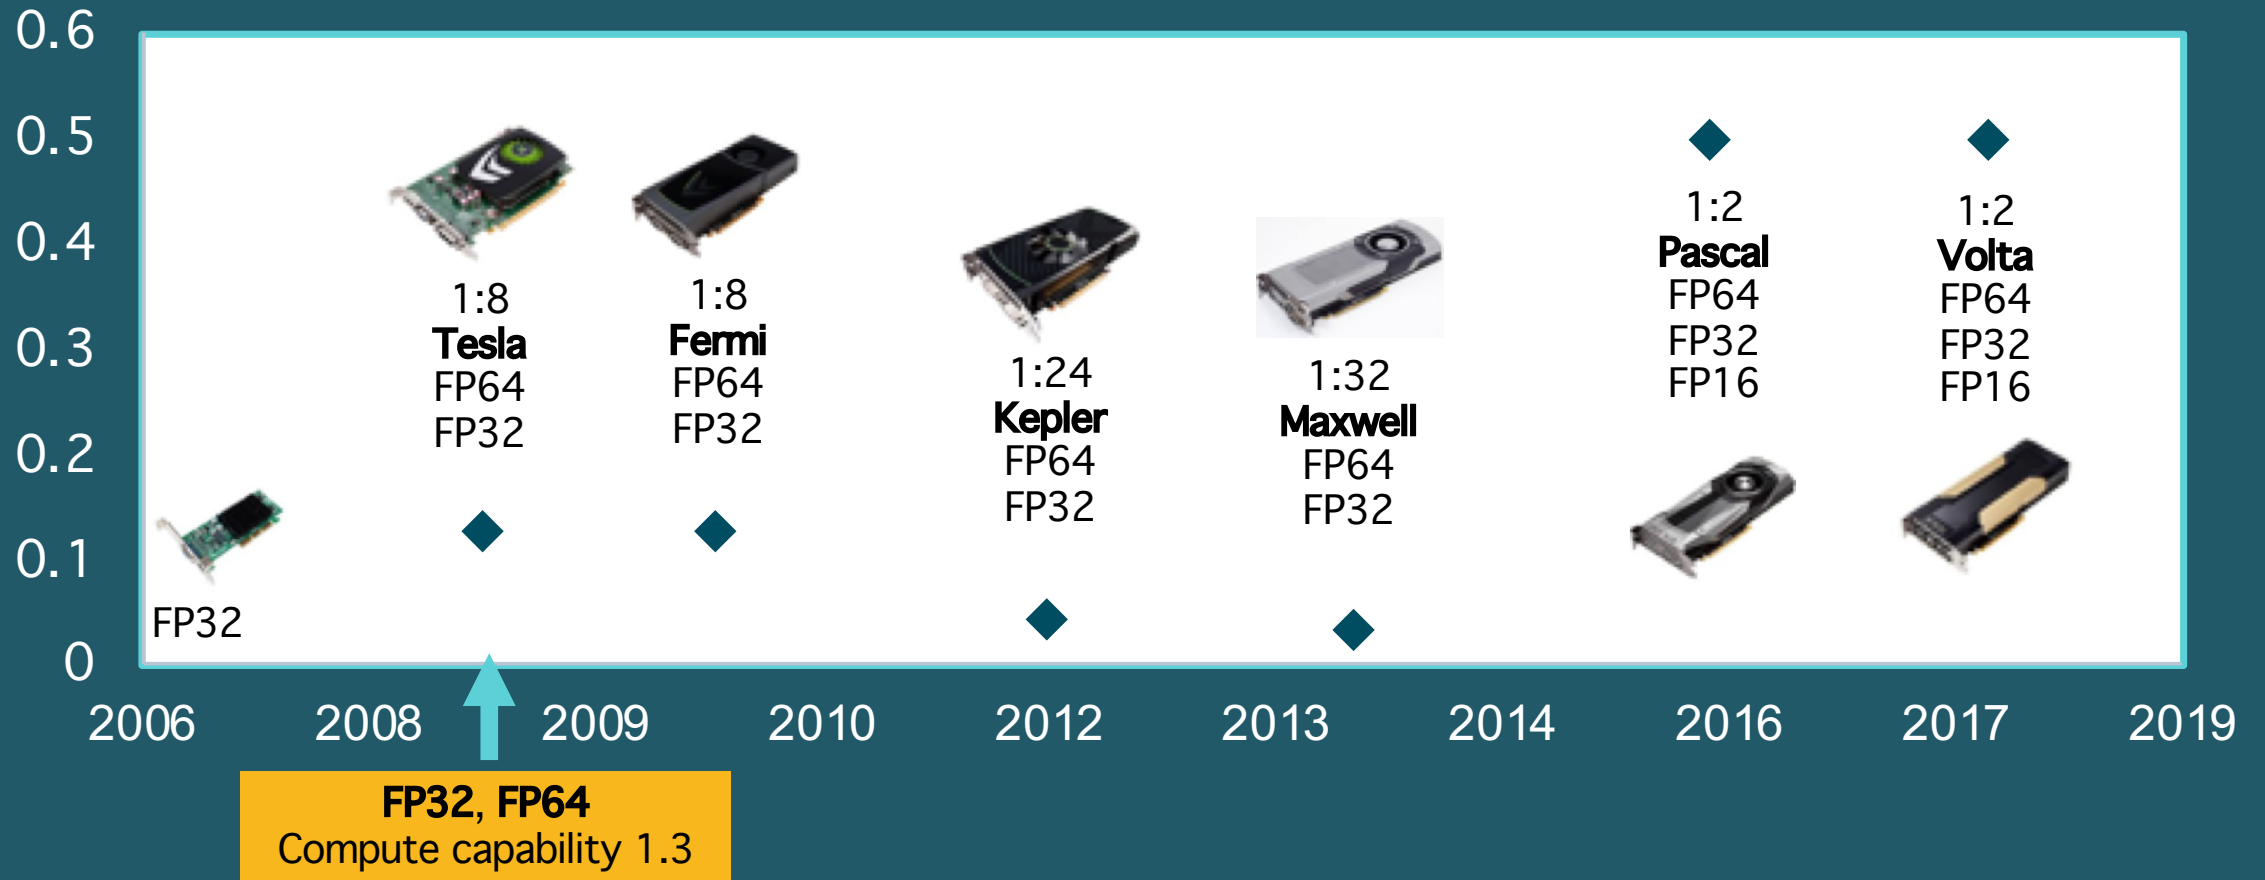
# FPChecker: Automatic Detection of Floating-Point Exceptions in GPUs



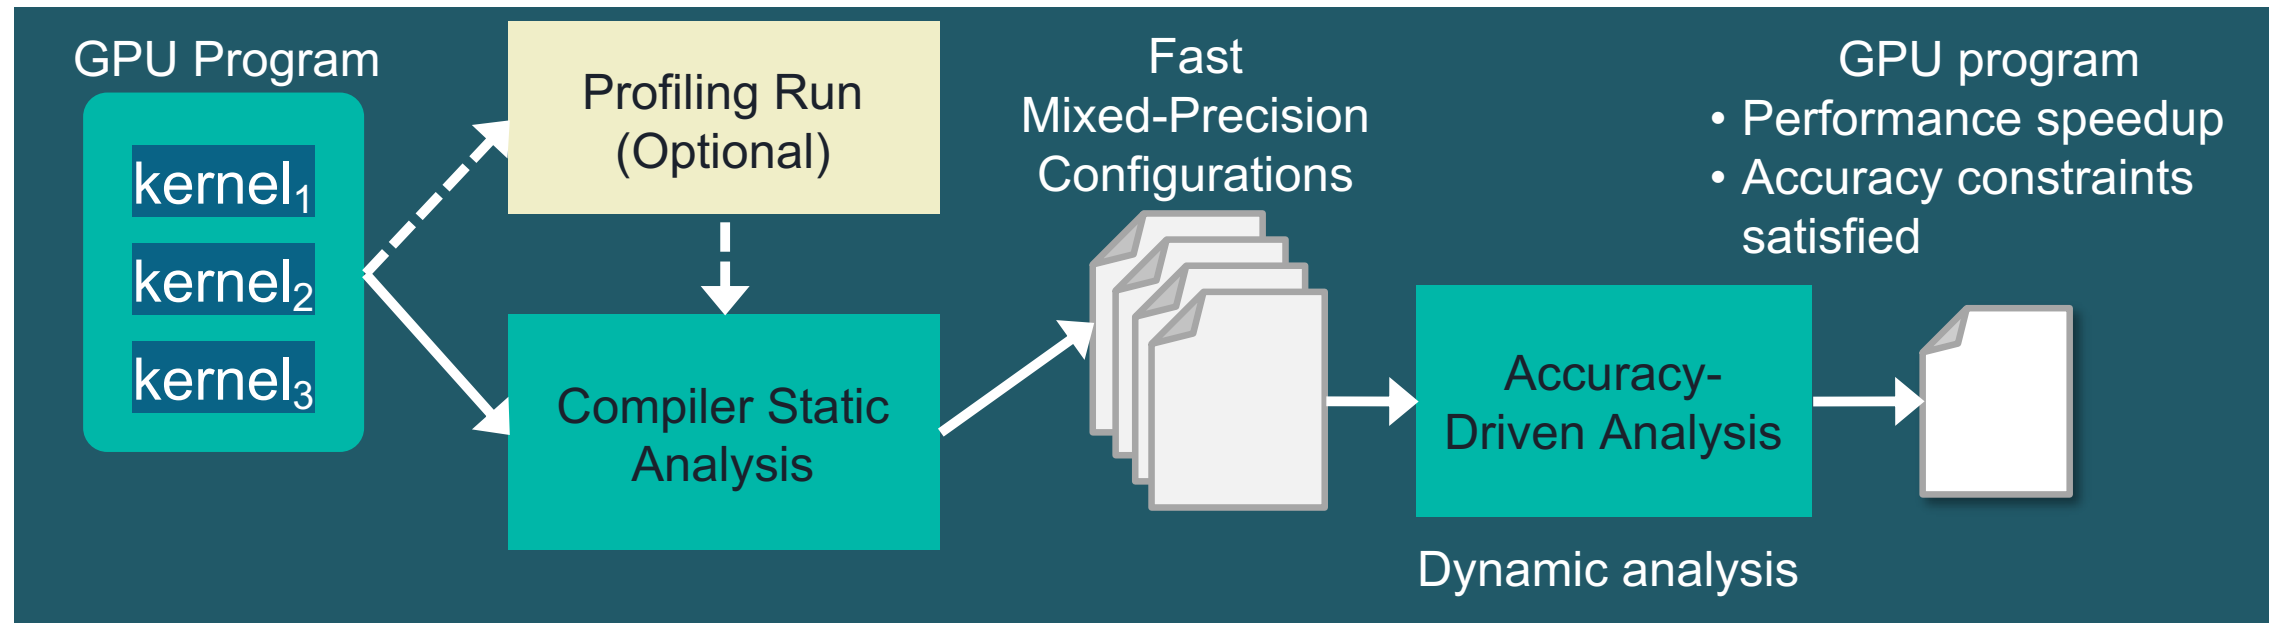
<https://github.com/LLNL/FPChecker>



# Floating-Point Precision Levels in GPUs Are Increasing



# GPUMixer: Performance-Driven Floating-Point Tuning for GPU Scientific Applications



Ignacio Laguna, Paul C. Wood, Ranvijay Singh, Saurabh Bagchi. **GPUMixer: Performance-Driven Floating-Point Tuning for GPU Scientific Applications**. ISC High Performance, Frankfurt, Germany, Jun 16-20, 2019 (**Best paper**)



# Tutorial on Floating-Point Analysis Tools

<http://fpanalysistools.org/>



- Demonstrate several analysis tools
- Hands-on exercises
- Cover various important aspects of floating-point and repro
- Tutorials:
  - LANL, Jan 9<sup>th</sup>, 2020
  - SC19, Denver, Nov 17<sup>th</sup>, 2019
  - PEARC19, Chicago, Jul 30<sup>th</sup>, 2019



#### Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.